

Flyto2 Core: Unauthenticated flyto-verification /run: callback_url SSRF and internal runner-secret exfiltration

Report Type: Weekly · Generated: July 30, 2026 17:59 UTC · Coverage: Jul 30 - Jul 30, 2026

1 Total Advisories	1 Critical	0 High	0 Medium
CRITICAL Severity	9.3 CVSS / Risk Score	TLP:CLEAR TLP Classification	PRO Customer Tier

Executive Summary

Summary The standalone `flyto-verification` service exposes `POST /run` with **no authentication**, on all interfaces (0.0.0.0:8344 per the shipped Dockerfile). The request body's `callback\_url` is used verbatim for an outbound POST that **unconditionally attaches** `X-Internal-Key: \$FLYTO\_RUNNER\_SECRET`. The `callback\_url` bypasses the service's `target\_allowed` allowlist (which only inspects `params.target\_url`) and is never passed through any SSRF guard. This yields (a) unauthenticated SSRF to internal/metadata endpoints with an attacker-controlled JSON body, and (b) exfiltration of the internal runner secret to an attacker-controlled host — allowing forged authenticated callbacks to the real engine. ## Root Cause - `src/core/verification\_service.py:363-364` — `/run` has no `Depends`/auth dependency. - `resolve\_callback\_url` returns the client `callback\_url` verbatim (:315-316). - `post\_callback` attaches `X-Internal-Key: \$FLYTO\_RUNNER\_SECRET` whenever the env var is set (:32

CONFIDENTIALITY NOTICE: This report is classified PRO tier and intended exclusively for the named recipient. Redistribution, resale, or disclosure to unauthorized parties is strictly prohibited. © 2026 CYBERDUDEBIVASH SENTINEL APEX. All rights reserved.

Top Threat Advisories

High-priority advisories ranked by CVSS score, exploitability, and actor attribution.

Advisory Title	Severity	CVSS	EPSS	AI Summary
Flyto2 Core: Unauthenticated flyto-verification /run: callBack SSR9.3	CRITICAL	SSR9.3	29.0%	## Summary The standalone `flyto-verification` service exposes `POST /run` wi

MITRE ATT&CK® v15 Mapping

Technique density score: 0.0 · Techniques observed: 0 · Tactics covered: 0

MITRE ATT&CK techniques pending analyst enrichment. Deploy detection rules from Section 6 to identify related activity.

Threat Actor Intelligence

Attribution analysis across advisories in this report period.

Threat Actor	Count	Associated CVEs / Indicators
CDB-UNATTR-CVE	1	-

Indicators of Compromise (IOCs)

IOC extraction for this advisory is pending enrichment pipeline completion. PRO subscribers receive real-time IOC push to SIEM/SOAR via the SENTINEL APEX webhook at: **GET /api/v1/intel/iocs?advisory={id}**. Full IOC packages include IPv4, domain, SHA256, and URL indicators with confidence scores, first-seen timestamps, and STIX 2.1 formatted bundles.

Detection Engineering

Deploy these production-ready detection rules into your SIEM platform to identify exploitation attempts in real time. Compatible with Microsoft Sentinel, Splunk ES, and any Sigma-compatible platform.

Sigma Rule - Webserver / Process Monitoring

```
title: SENTINEL APEX - CDB-ADVISORY Exploitation Attempt
id: cdb-cdb-advisory-det
status: experimental
description: Detects exploitation indicators related to CDB-ADVISORY.
references:
- https://intel.cyberdudebivash.com
author: CyberDudeBivash SENTINEL APEX
date: 2026/07/30
tags:
- attack.initial_access
- attack.t1190
logsource:
category: webserver
detection:
keywords:
- 'CDB-ADVISORY'
condition: keywords
falsepositives:
- Security testing / penetration testing
level: medium
```

Microsoft Sentinel - KQL Query

```
// SENTINEL APEX - Flyto2 Core: Unauthenticated flyto-verification /run: callback_url SSRF and internal runner-secret exfiltration
// Advisory | Microsoft Sentinel KQL
let time_window = ago(24h);
SecurityAlert
| where TimeGenerated > time_window
| where Description has "Flyto2 Core: Unauthenticated flyto-verif"
or AlertName has "Flyto2 Core: Unauthenticated flyto-verif"
or ExtendedProperties has "Flyto2 Core: Unauthenticated flyto-verif"
| extend Rule = "APEX-ADVISORY", Severity = "CRITICAL"
| project TimeGenerated, AlertName, AlertSeverity, Entities, Description, Rule, Severity
| order by TimeGenerated desc
```

Incident Response Playbook

Phase 1 - 0-24 Hours (Containment)

1. Activate IR team; assign lead analyst and executive sponsor.
2. Isolate or WAF-shield Flyto2 Core: Unauthenticated flyto instances exposed to the internet.
3. Enable verbose access logging on all Flyto2 Core: Unauthenticated flyto endpoints immediately.

- 4. Pull last 72h of web server and application logs for forensic baselining.
- 5. Check SENTINEL APEX IOC feeds and SIEM alerts for exploitation hits in your environment.
- 6. Notify relevant internal stakeholders and legal / compliance team.

Phase 2 - 24-72 Hours (Eradication)

- 1. Apply vendor patch for Flyto2 Core: Unauthenticated flyto or implement recommended mitigation.
- 2. Deploy Sigma and KQL detection rules from Section 6 across all SIEM platforms.
- 3. Rotate all API keys, service account credentials, and session tokens.
- 4. Conduct forensic timeline reconstruction for potential compromise window.
- 5. Validate patch deployment across 100% of affected instances via asset inventory.

Phase 3 - 7 Days (Recovery & Hardening)

- 1. Conduct post-incident review and update incident response runbooks.
- 2. Threat hunt for lateral movement or persistence artifacts using MITRE ATT&CK mapping.
- 3. Update asset inventory with patched version information and closure evidence.
- 4. Submit findings to internal risk register and CISO dashboard.
- 5. Schedule follow-up vulnerability scan to confirm remediation completeness.

Financial Impact Analysis (FAIR Model)

\$2.4M - \$12M Breach Cost Range (FAIR)	\$5.8M IBM Cost of Breach 2025 Median	\$1.2M/day Business Interruption Cost	\$850K Regulatory Fine Exposure
--	--	--	--

Critical severity breaches trigger mandatory 72-hour regulatory disclosure (GDPR Art. 33), class-action exposure, and long-tail legal costs averaging 18-36 months post-incident.

Remediation Cost Estimate: \$380K · Includes engineering time, IR retainer activation, forensic investigation, and customer notification costs.

Regulatory & Compliance Implications

Organizations processing this advisory must evaluate obligations under applicable regulatory frameworks. The table below maps this threat to relevant compliance requirements.

Framework	Reference	Obligation
GDPR	Art. 33 / Art. 83(4)	72-hour breach notification to DPA; fines up to 4% of global annual turnover
PCI-DSS v4.0	Req. 6.3 / 6.4	Critical patches within 1 month; compensating controls required
NIS2 Directive	Art. 23	Early warning within 24h; full incident report within 72h

Framework	Reference	Obligation
SOX / ICFR	Section 302 / 404	Material cybersecurity weakness disclosure in 10-K/10-Q
ISO 27001:2022	A.8.8 / A.5.29	Documented vulnerability management and incident management

Actionable Recommendations

Prioritized defensive actions based on this period's intelligence.

1. Apply patches immediately for Flyto2 Core: Unauthenticated flyto-verification /run: callback_url SSRF and internal runner-secret exfiltration.
2. Enable enhanced monitoring for related IOCs.
3. Review SIEM rules for associated MITRE ATT&CK techniques.
4. Apply vendor patch for Flyto2 Core: Unauthenticated flyto-verification /run: callback_url SSRF and internal runner-secret exfiltration immediately - check NVD for official fix guidance.
5. Enable enhanced logging and alerting on all systems running affected software versions.
6. Deploy the Sigma detection rule from this report into your SIEM platform (Sentinel / Splunk / Elastic).
7. Audit all internet-facing instances of affected products and confirm patch deployment.

Appendix - Report Metadata & Legal

Field	Value
Report ID	intel--d32178a2c0e3b706978655cb
Report Type	Weekly
Platform	CYBERDUDEBIVASH SENTINEL APEX v166.0 GOD-MODE
Generated At	2026-07-30T17:59:04.32921
Customer Tier	PRO
Customer Email	-
TLP	TLP: CLEAR
STIX Version	2.1
ATT&CK Version	v15

Legal & Usage Notice

This threat intelligence report is produced by CYBERDUDEBIVASH SENTINEL APEX, operated by Bivash Nath (GSTIN: 21ARKPN8270G1ZP). All intelligence is aggregated from public vulnerability databases (NVD, CVSS), threat feeds, and proprietary AI enrichment pipelines. This report is provided AS-IS for informational and defensive security purposes only. CYBERDUDEBIVASH accepts no liability for actions taken based on this intelligence. Redistribution, resale, or disclosure to unauthorized parties is strictly prohibited without written consent. For licensing inquiries: enterprise@cyberdudebivash.com